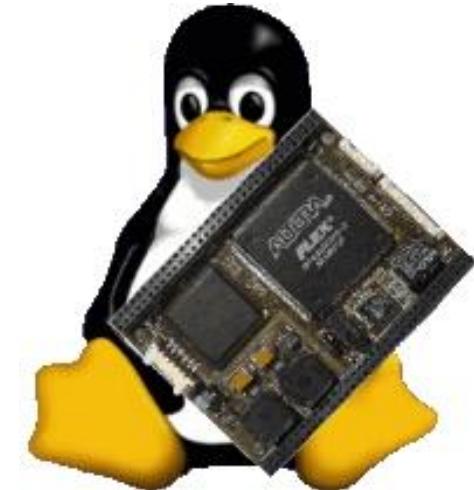
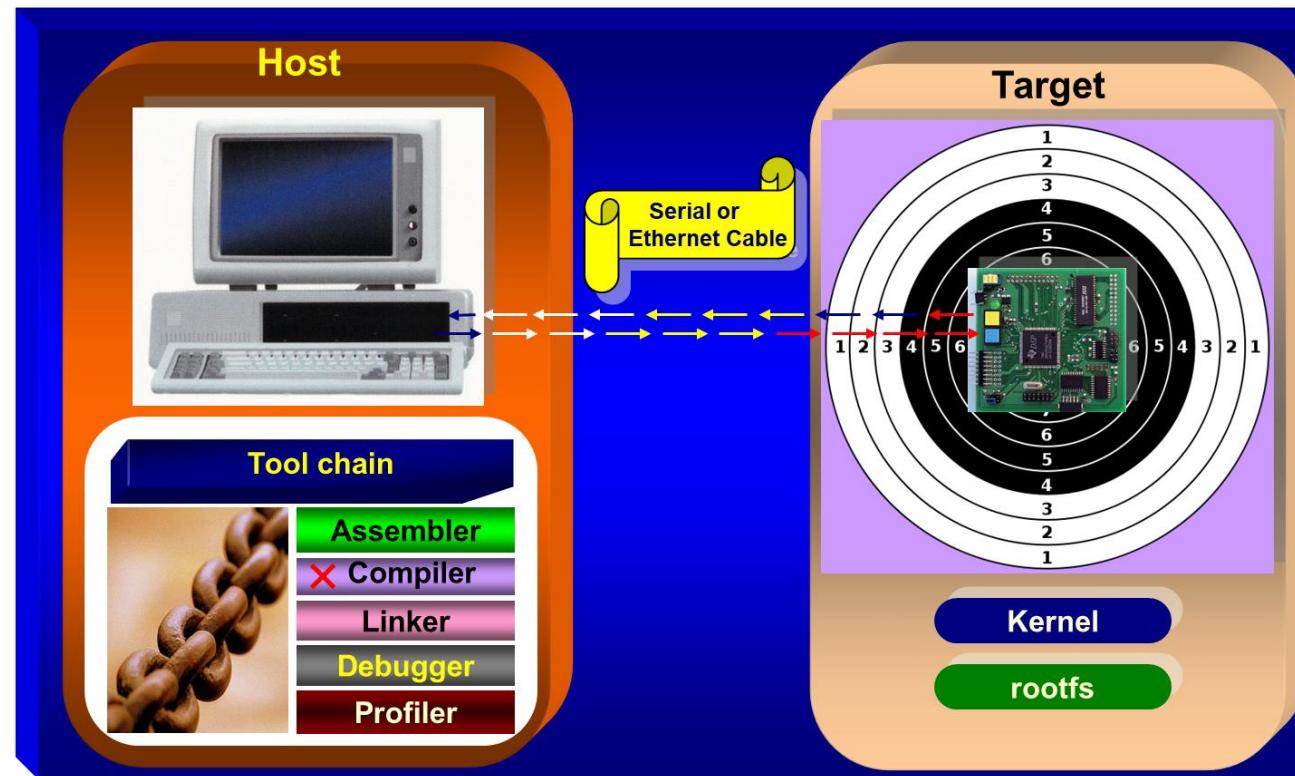
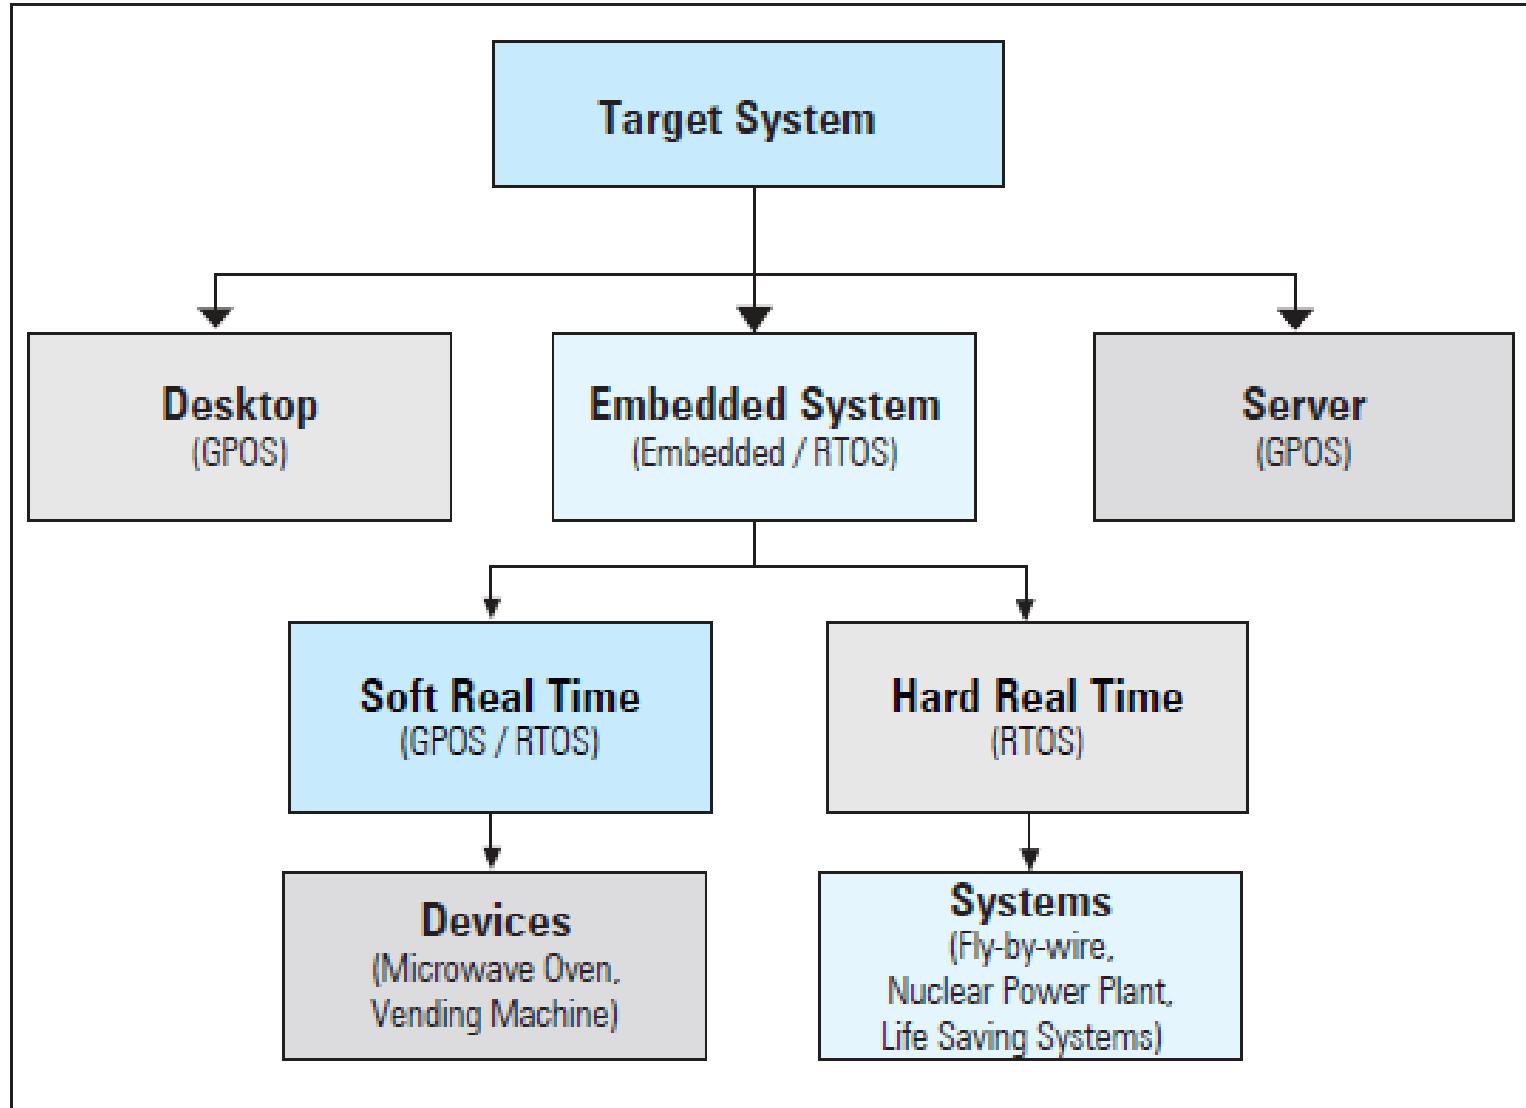


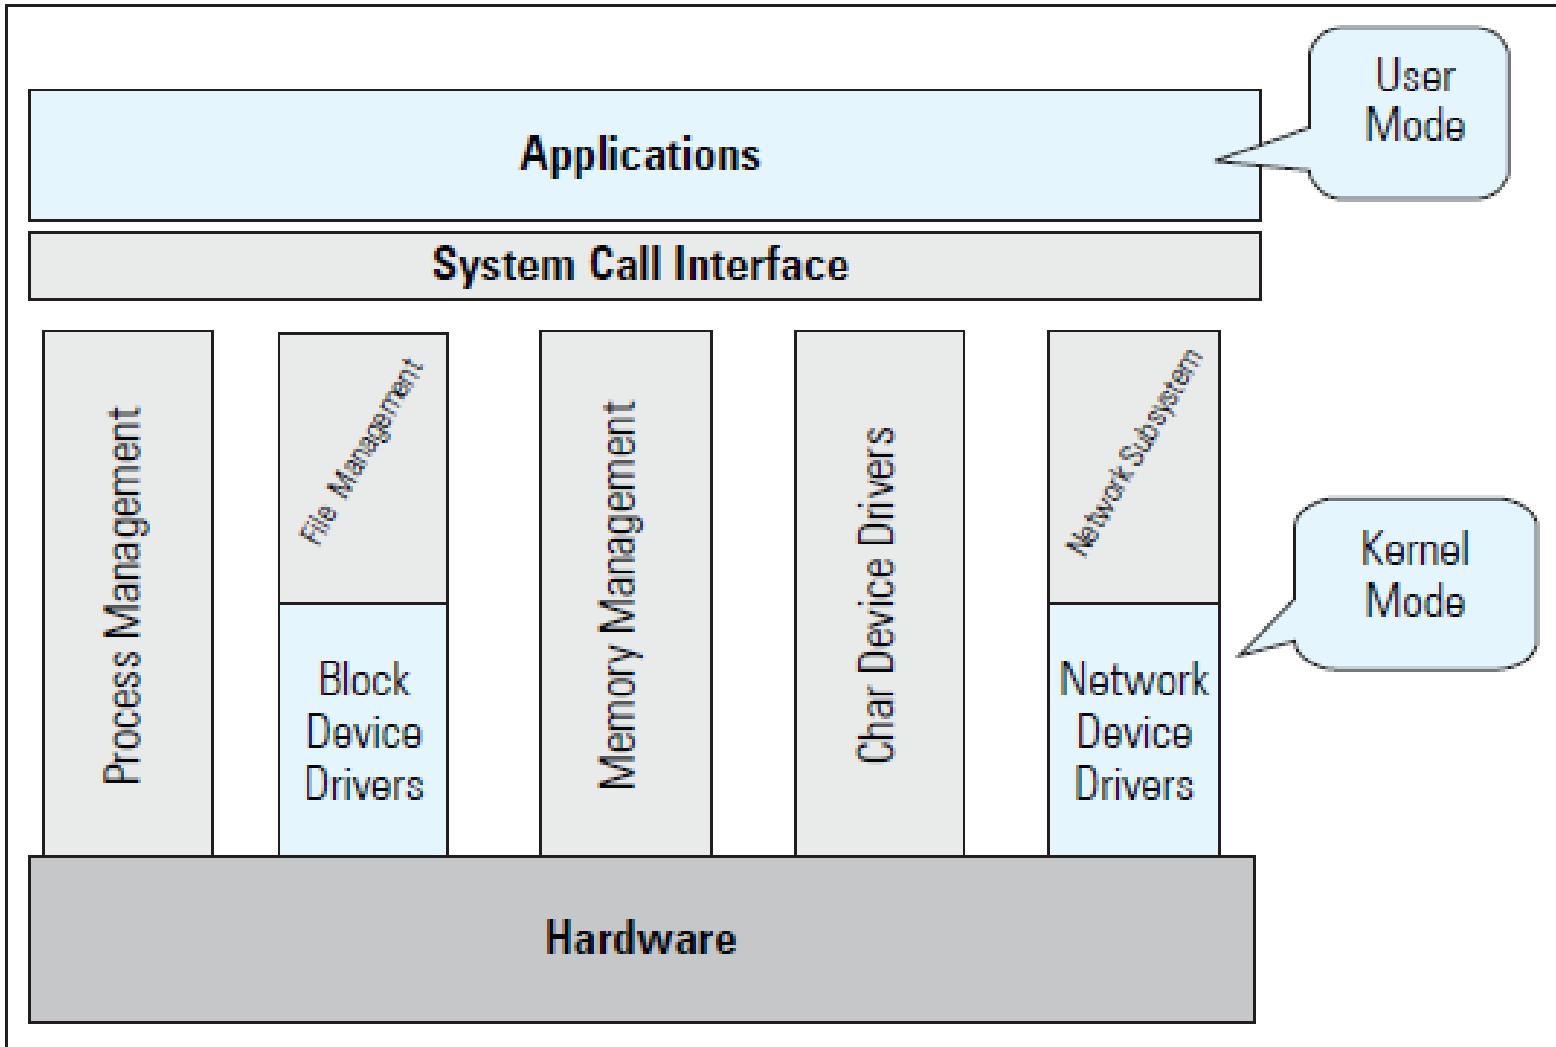
# CS 513 Software System Recap



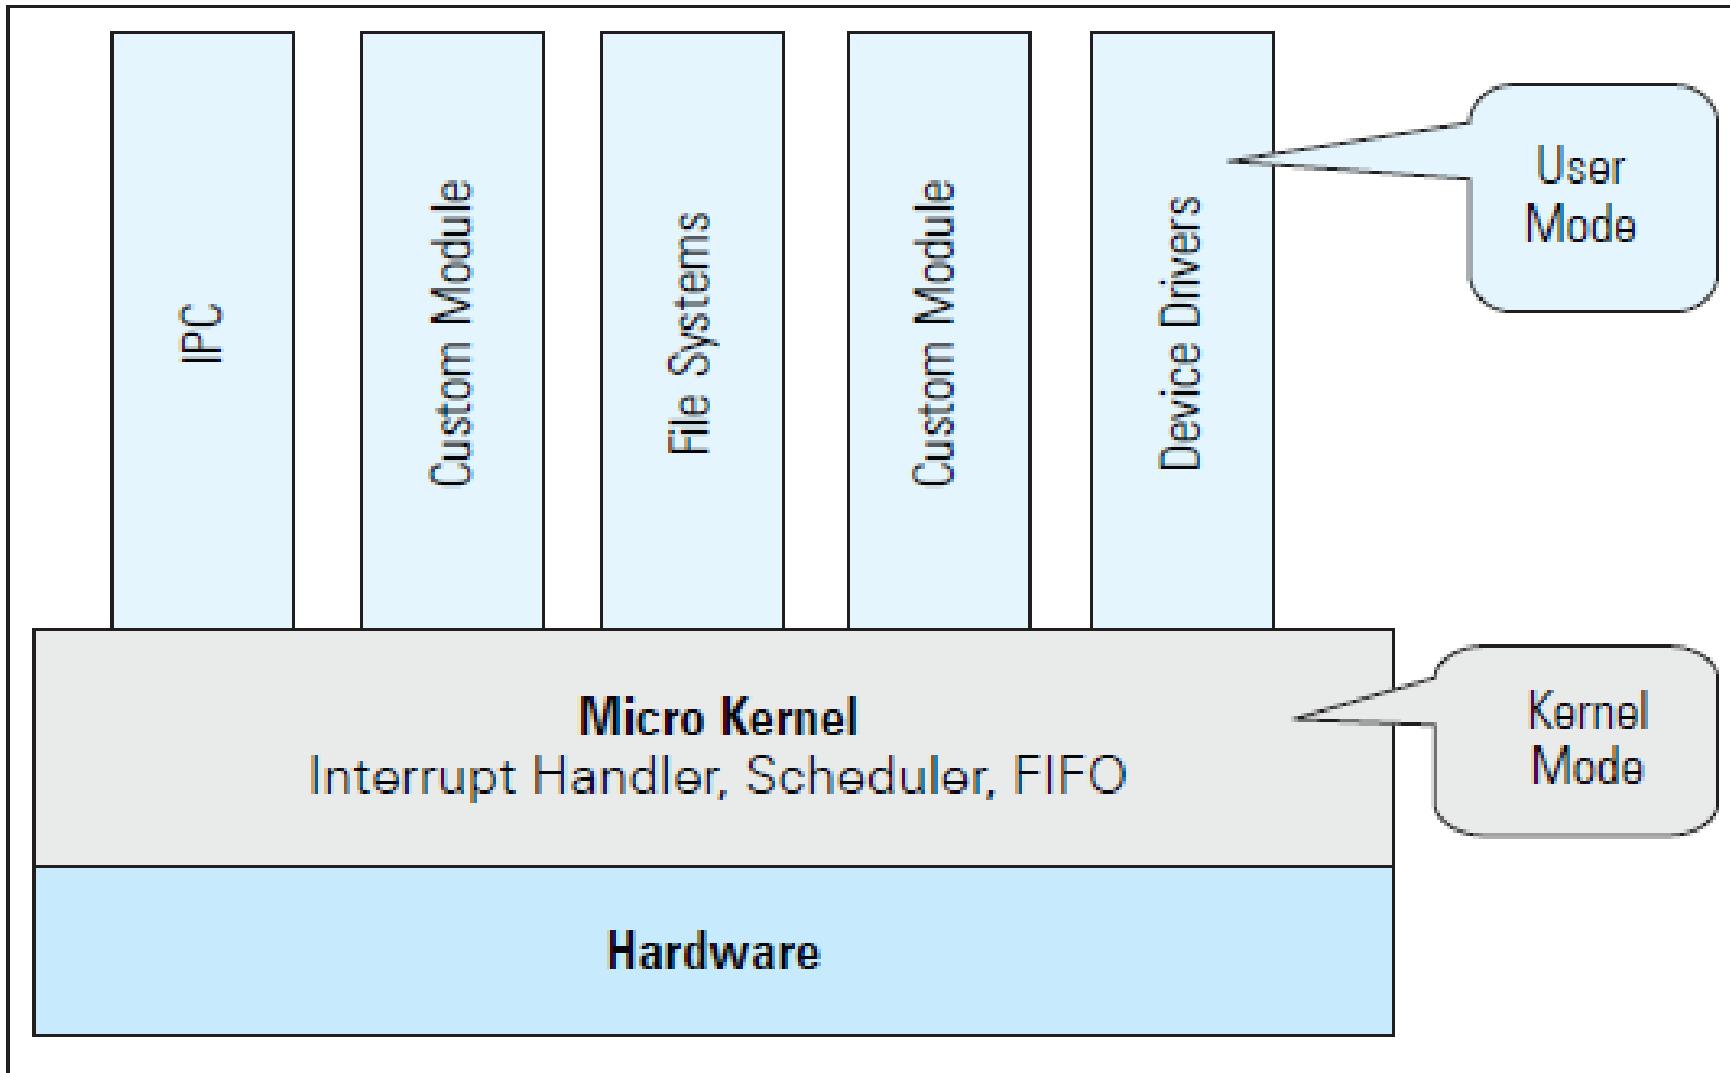
Prof. B Thangaraju

# Types of System

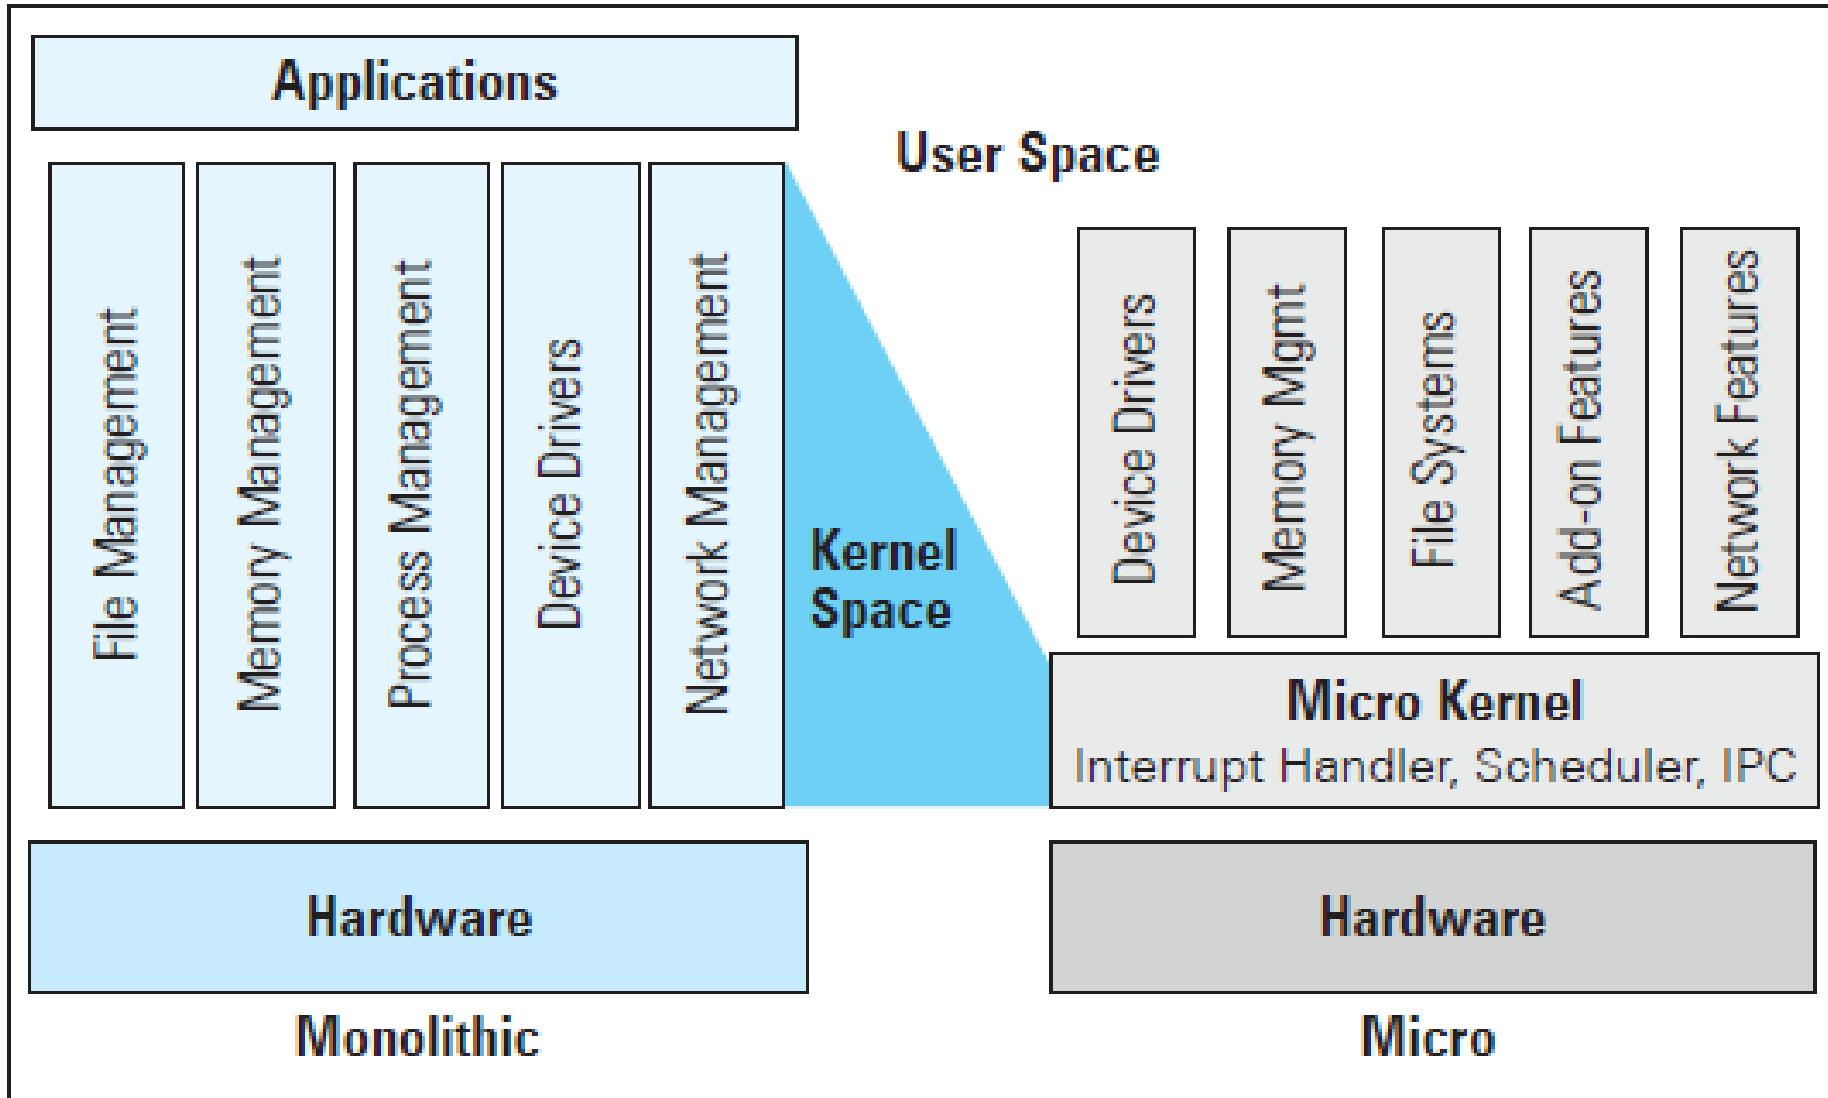




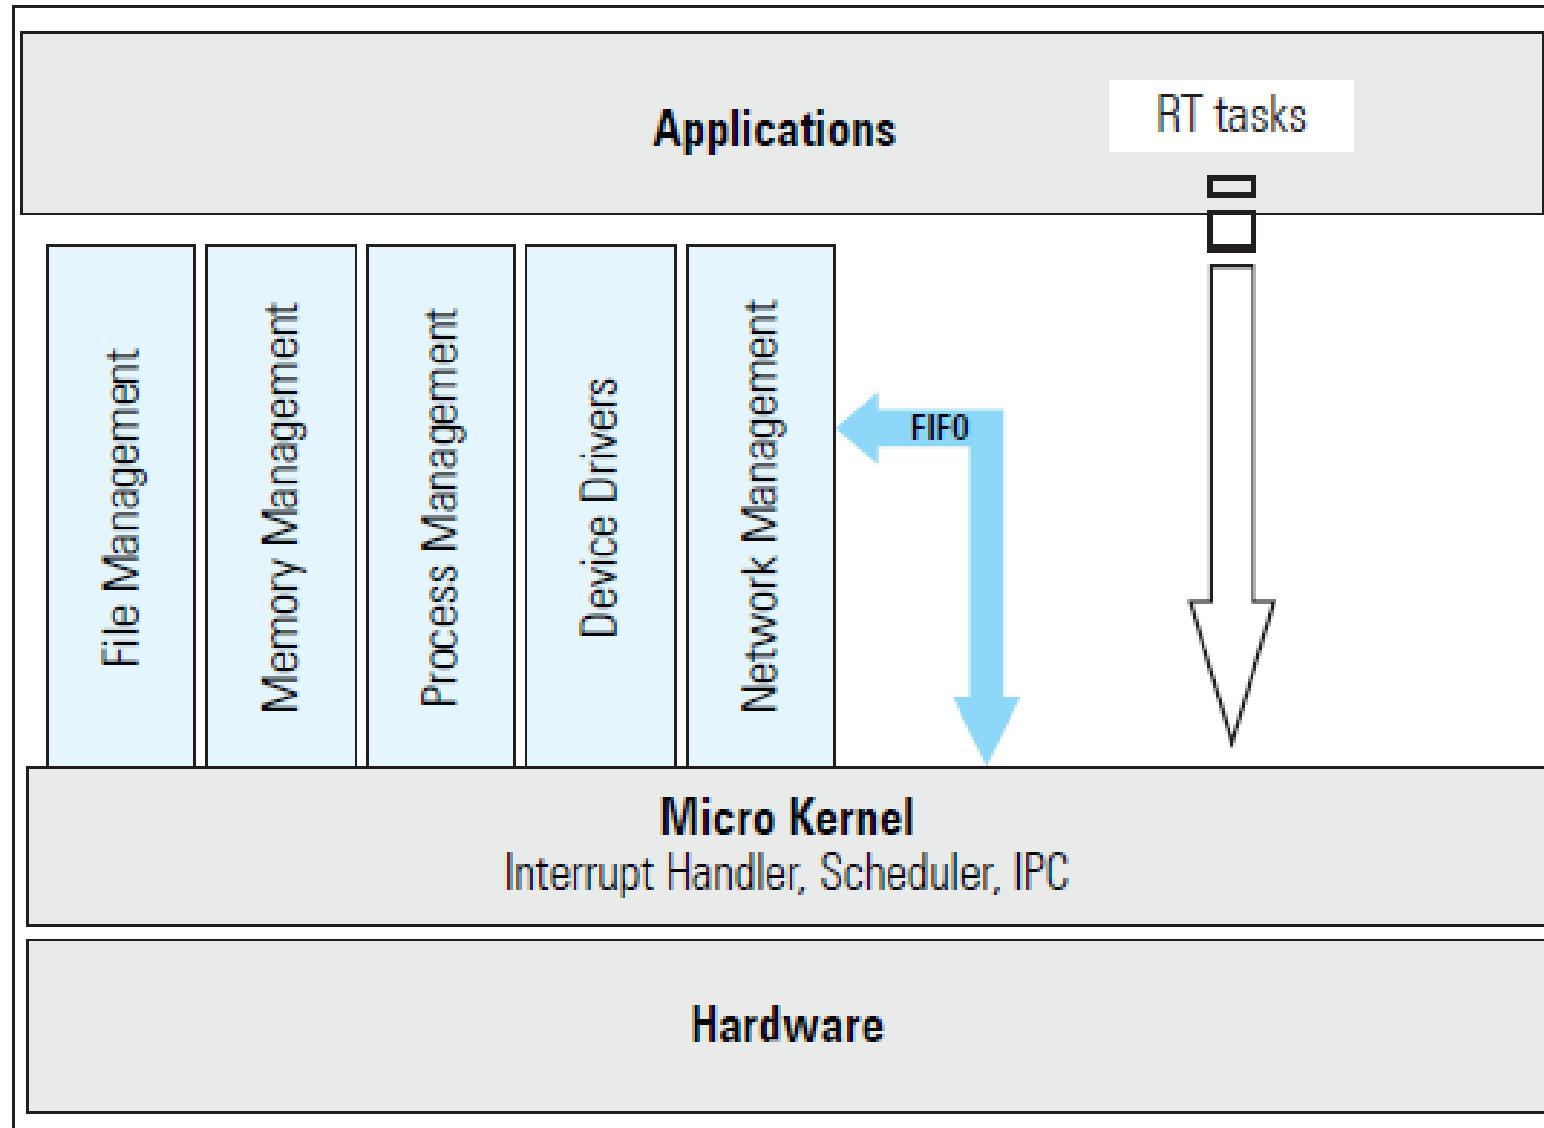
# Micro Kernel Architecture



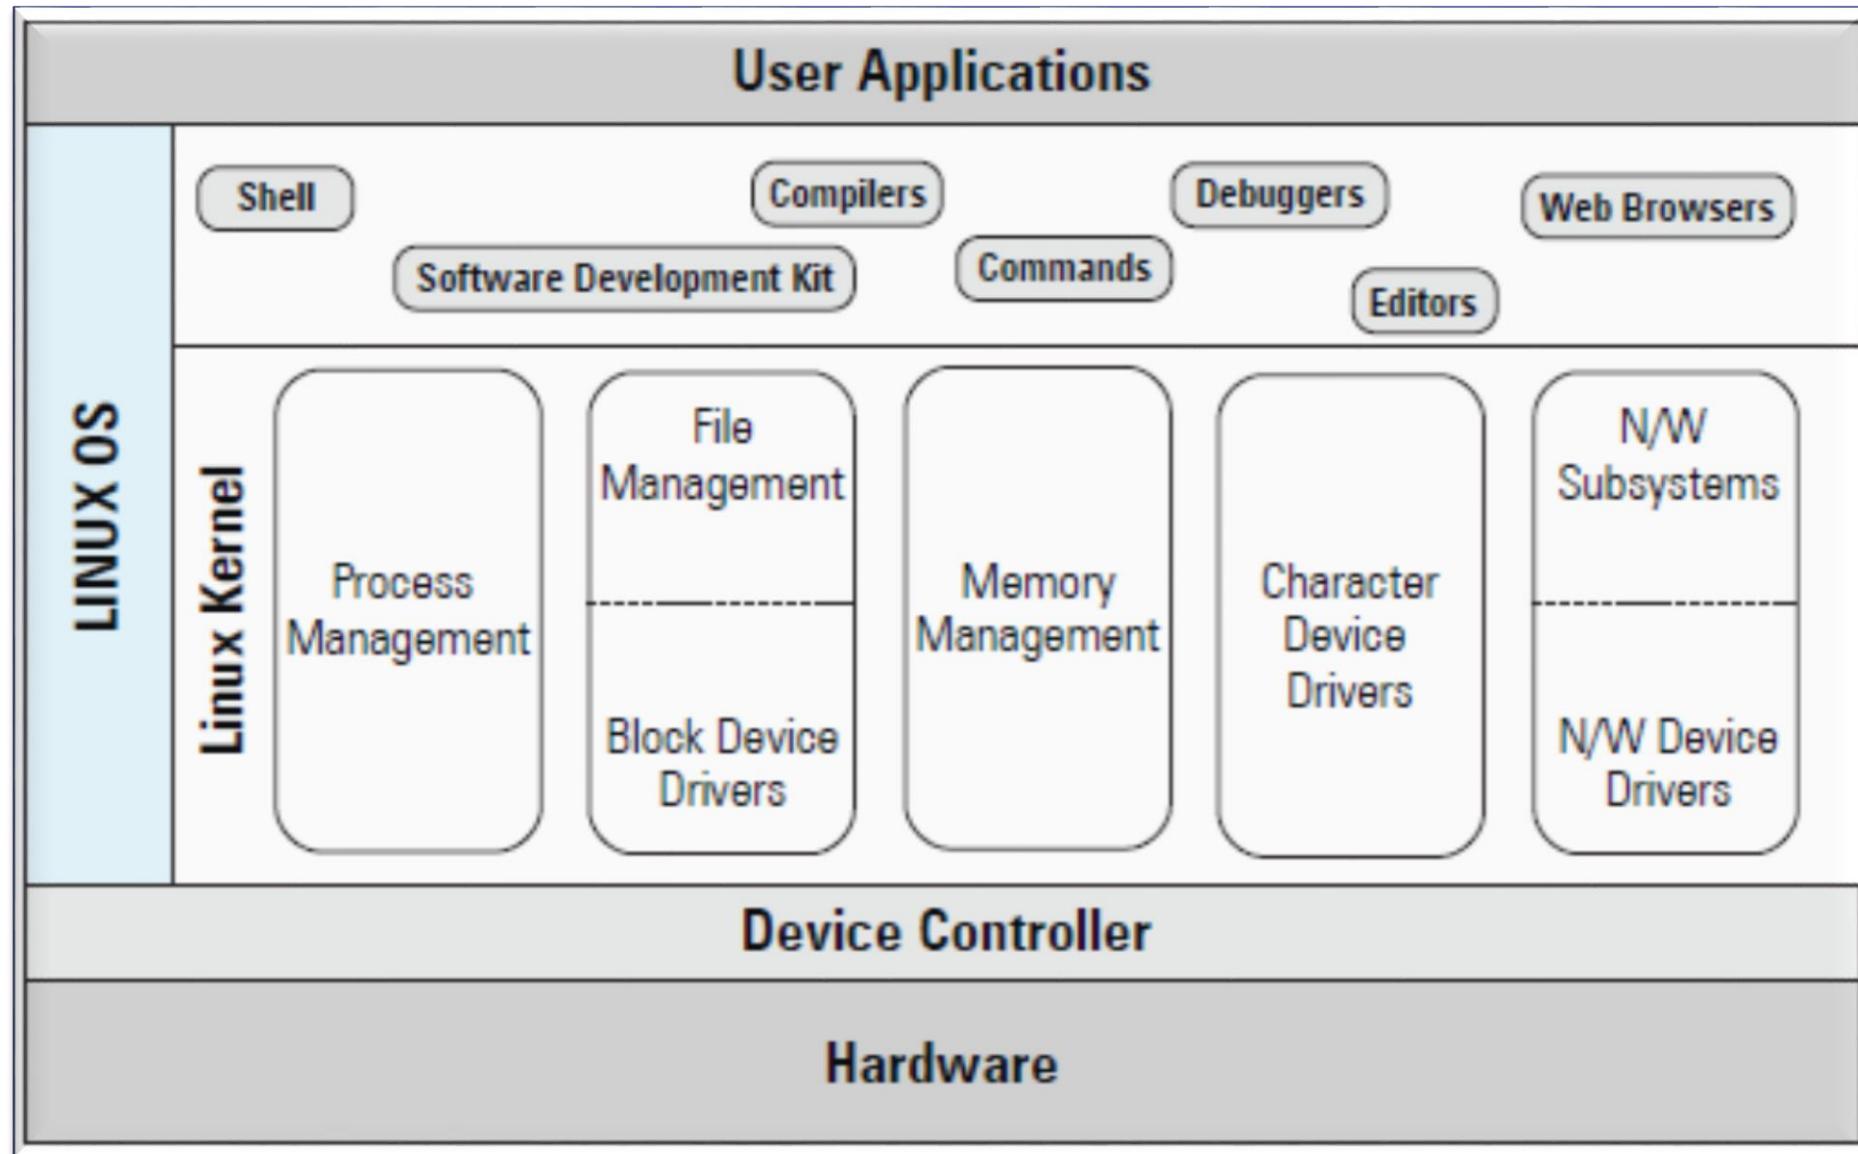
# Monolithic Vs Micro Kernel



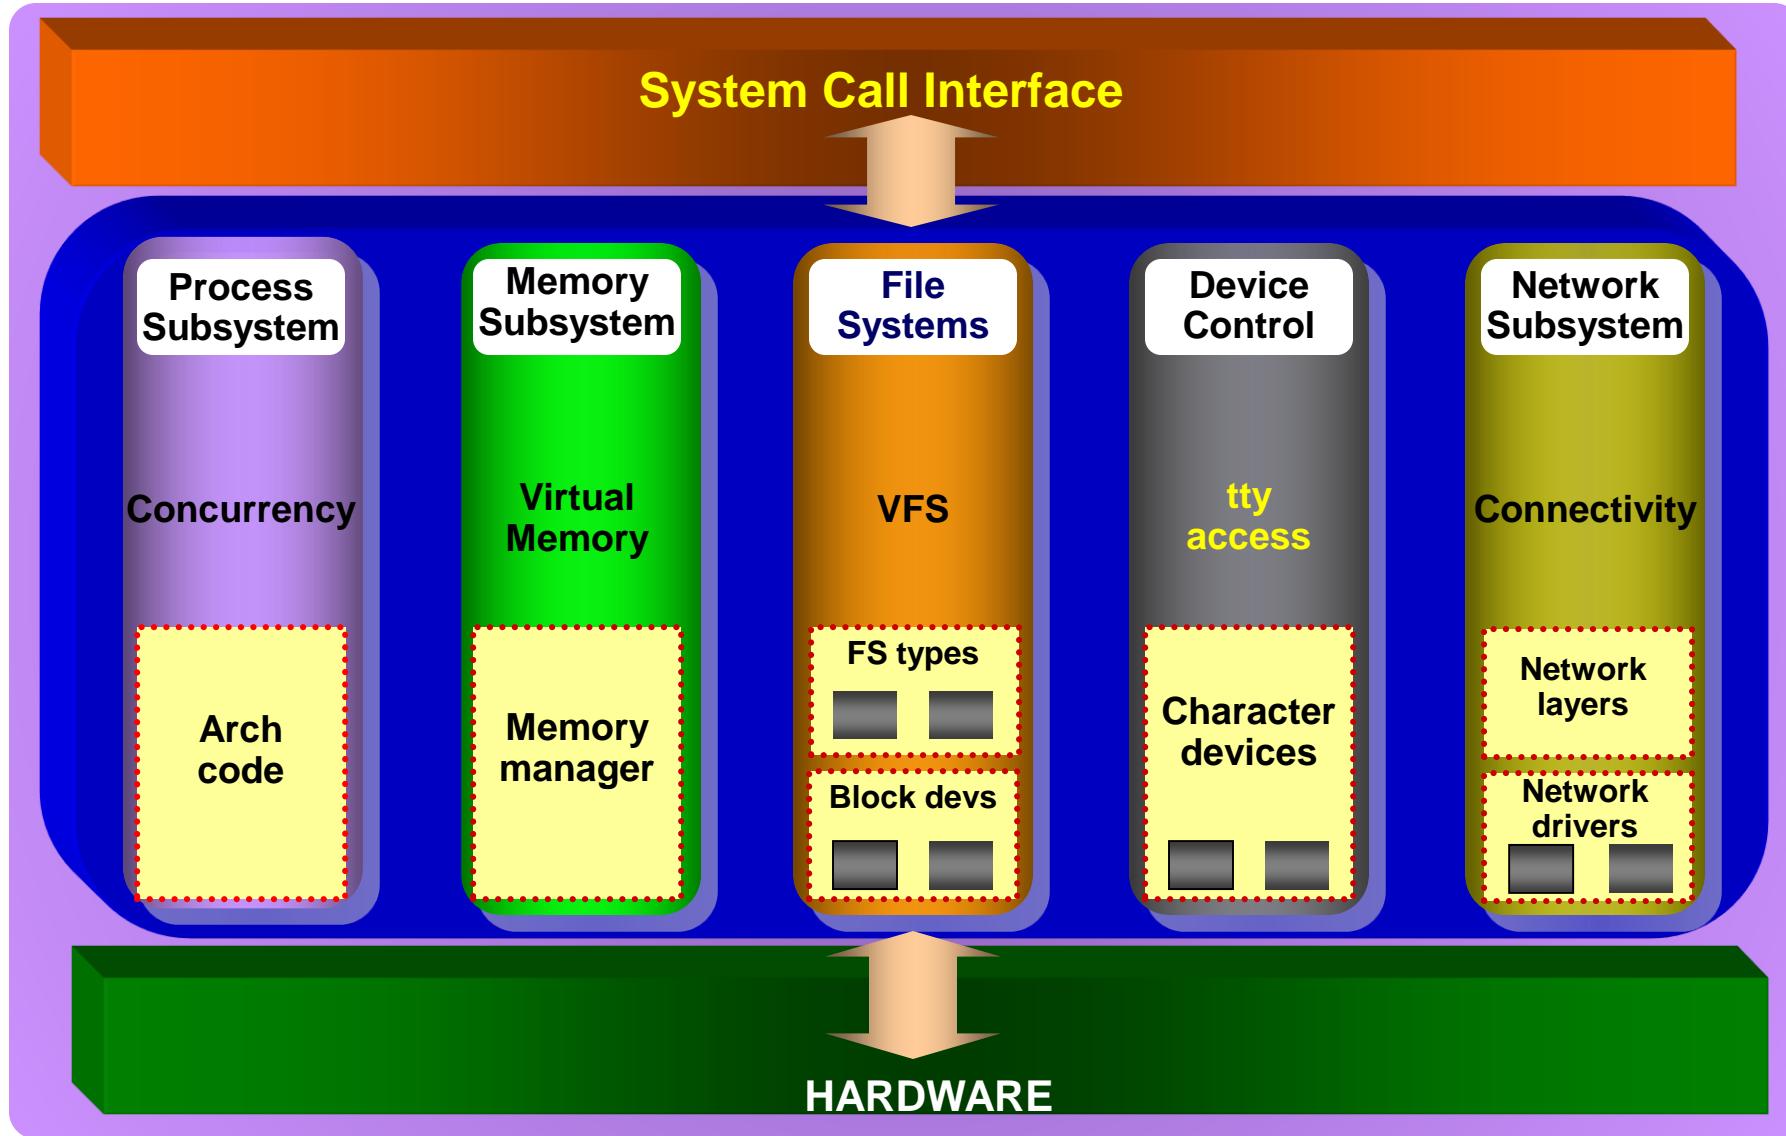
# Hybrid Kernel Architecture



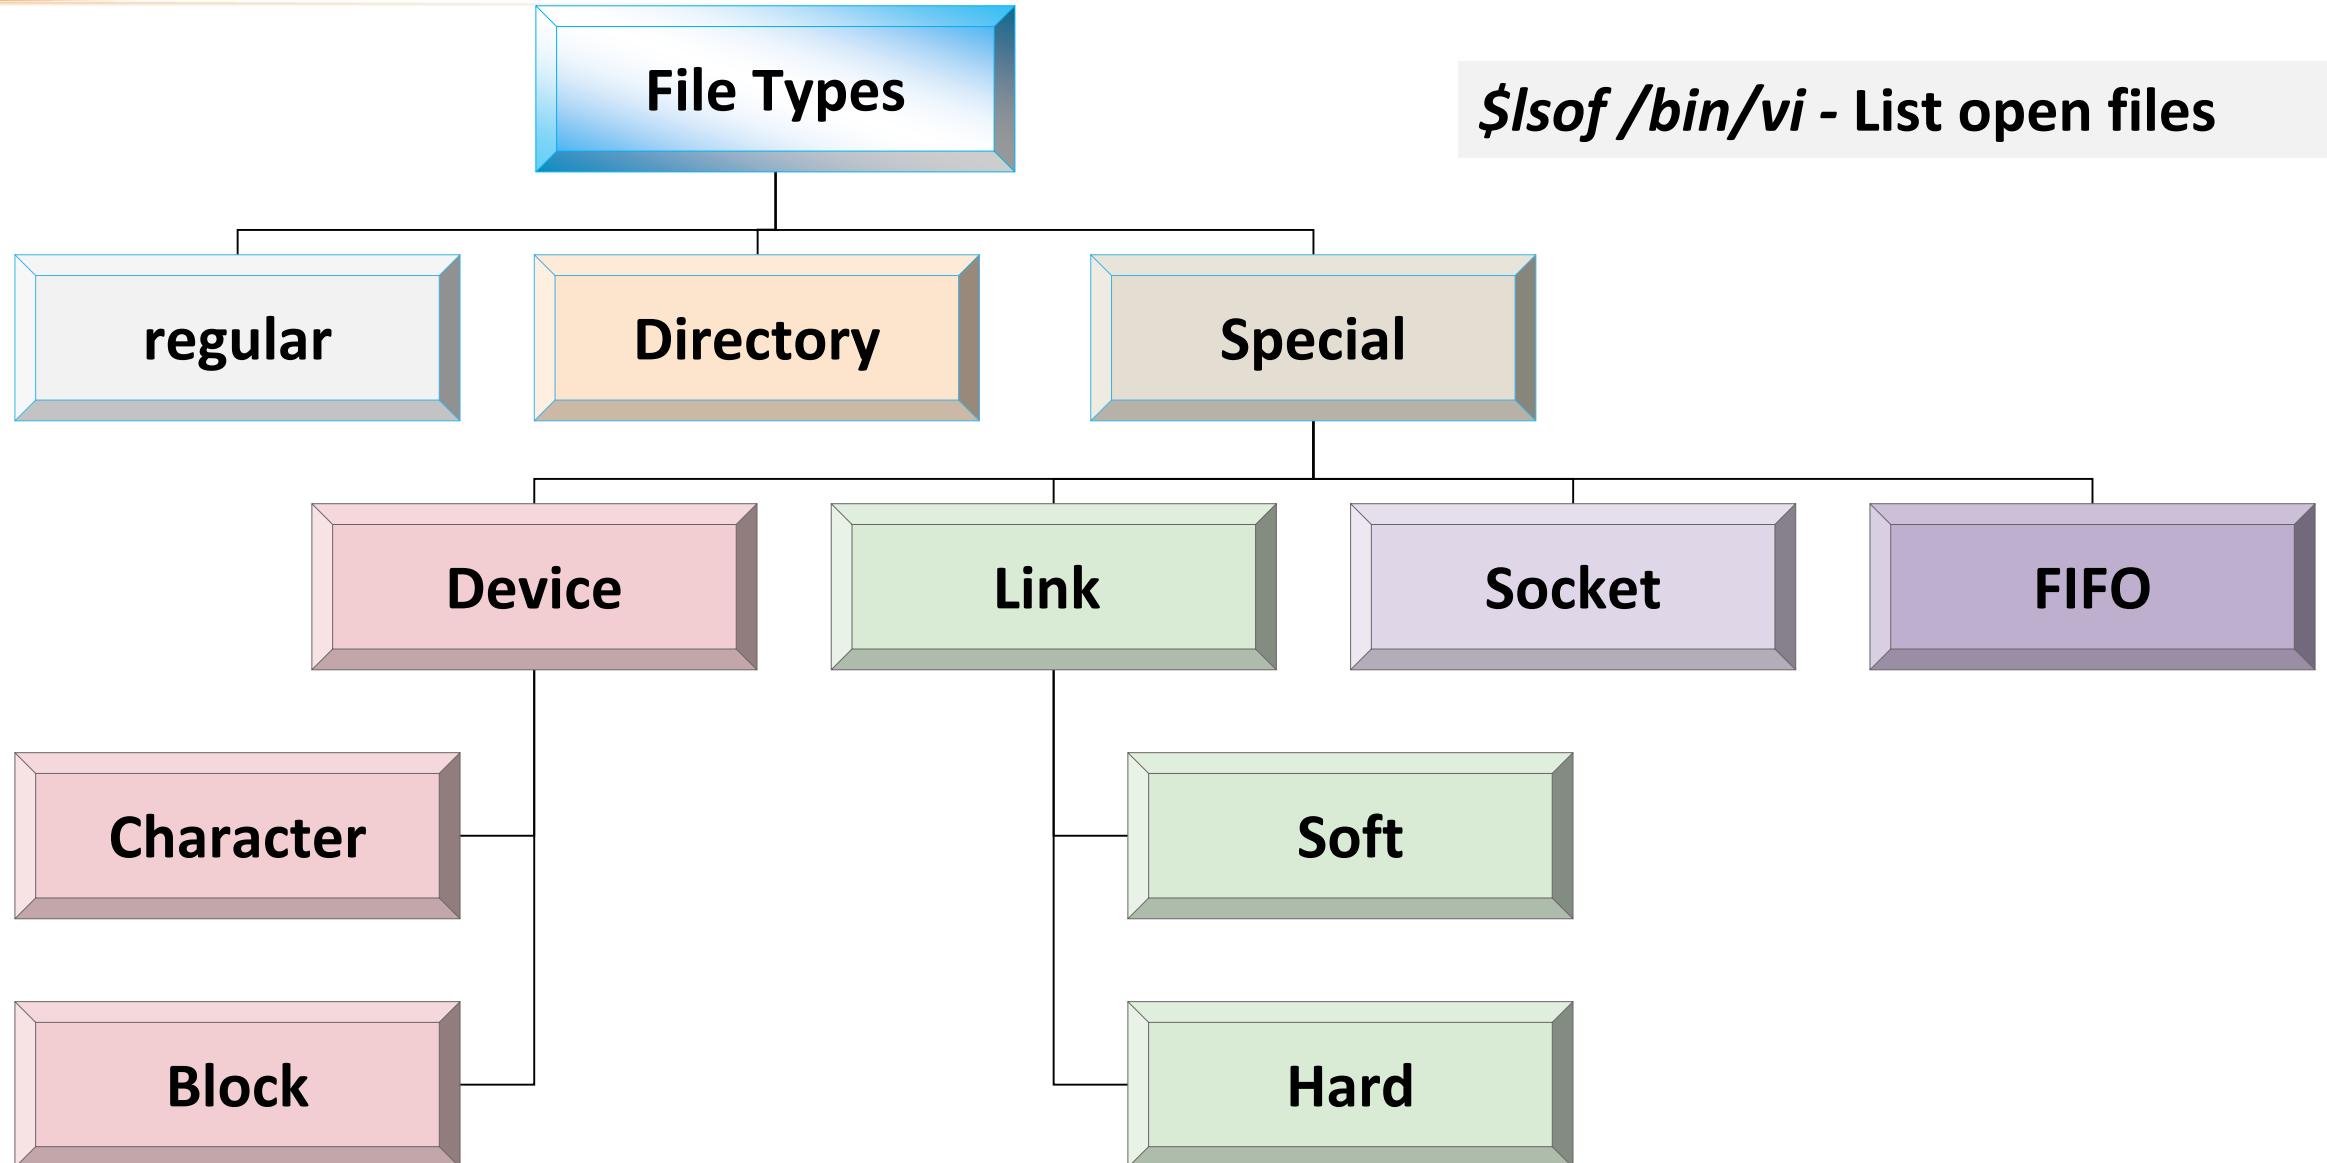
# Introduction – Linux OS



# Linux Kernel Architecture



# File Types



- Identification of File types
  - ? r w - r- - r- -
  - ?-specifies a type of a file
- Regular (-)
- Directory (d)
- Special Files

## Special Files

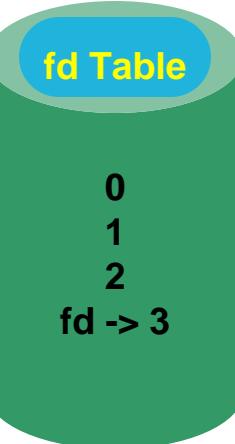
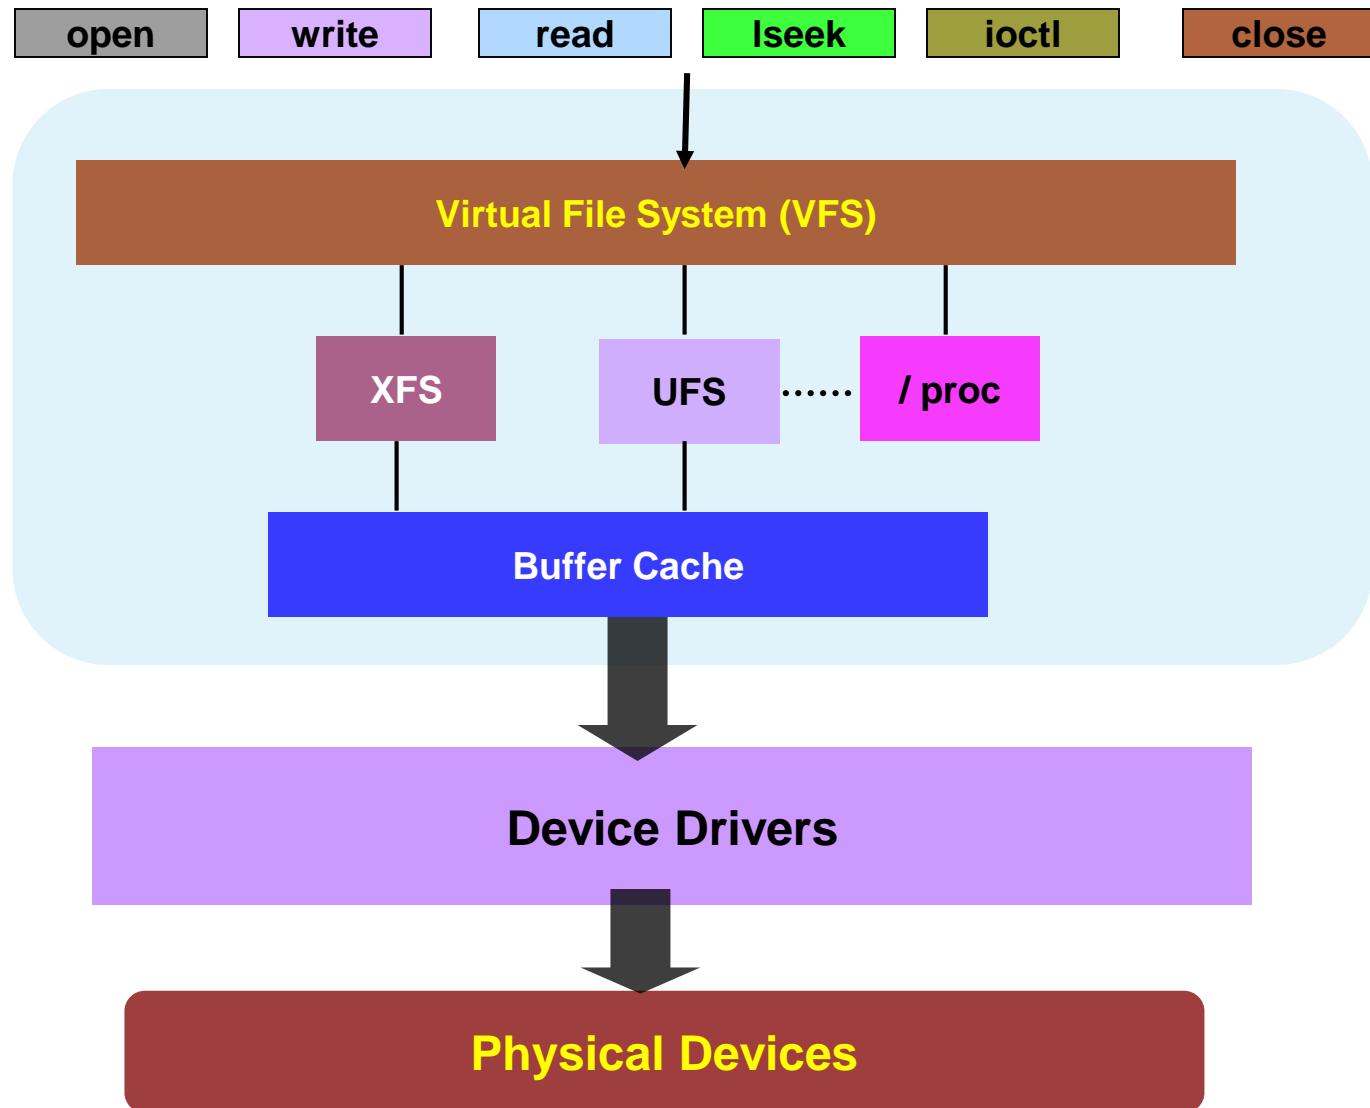
- FIFO (p)
- Socket (S)
- Link File
  - Soft Link (l)
  - Hard Link (inode numbers are same)
- Device File
  - Character (c)
    - Example: Monitor, Keyboard, Mouse, Tape
  - Block (b)
    - Example: Hard disk, CDROM, Floppy

# File System

- Facilitates persistent storage and data management
- Facilitates file related system calls.
- Different types of file system for different needs –depends on implementation
  - A logical file system appears as a single entity to the user process, but it may be composed of a number of physical file systems.

- VFS
- XFS
- ext4
- UFS
- proc
- msdos
- iso9660
- Vfat
- Aufs
- .....

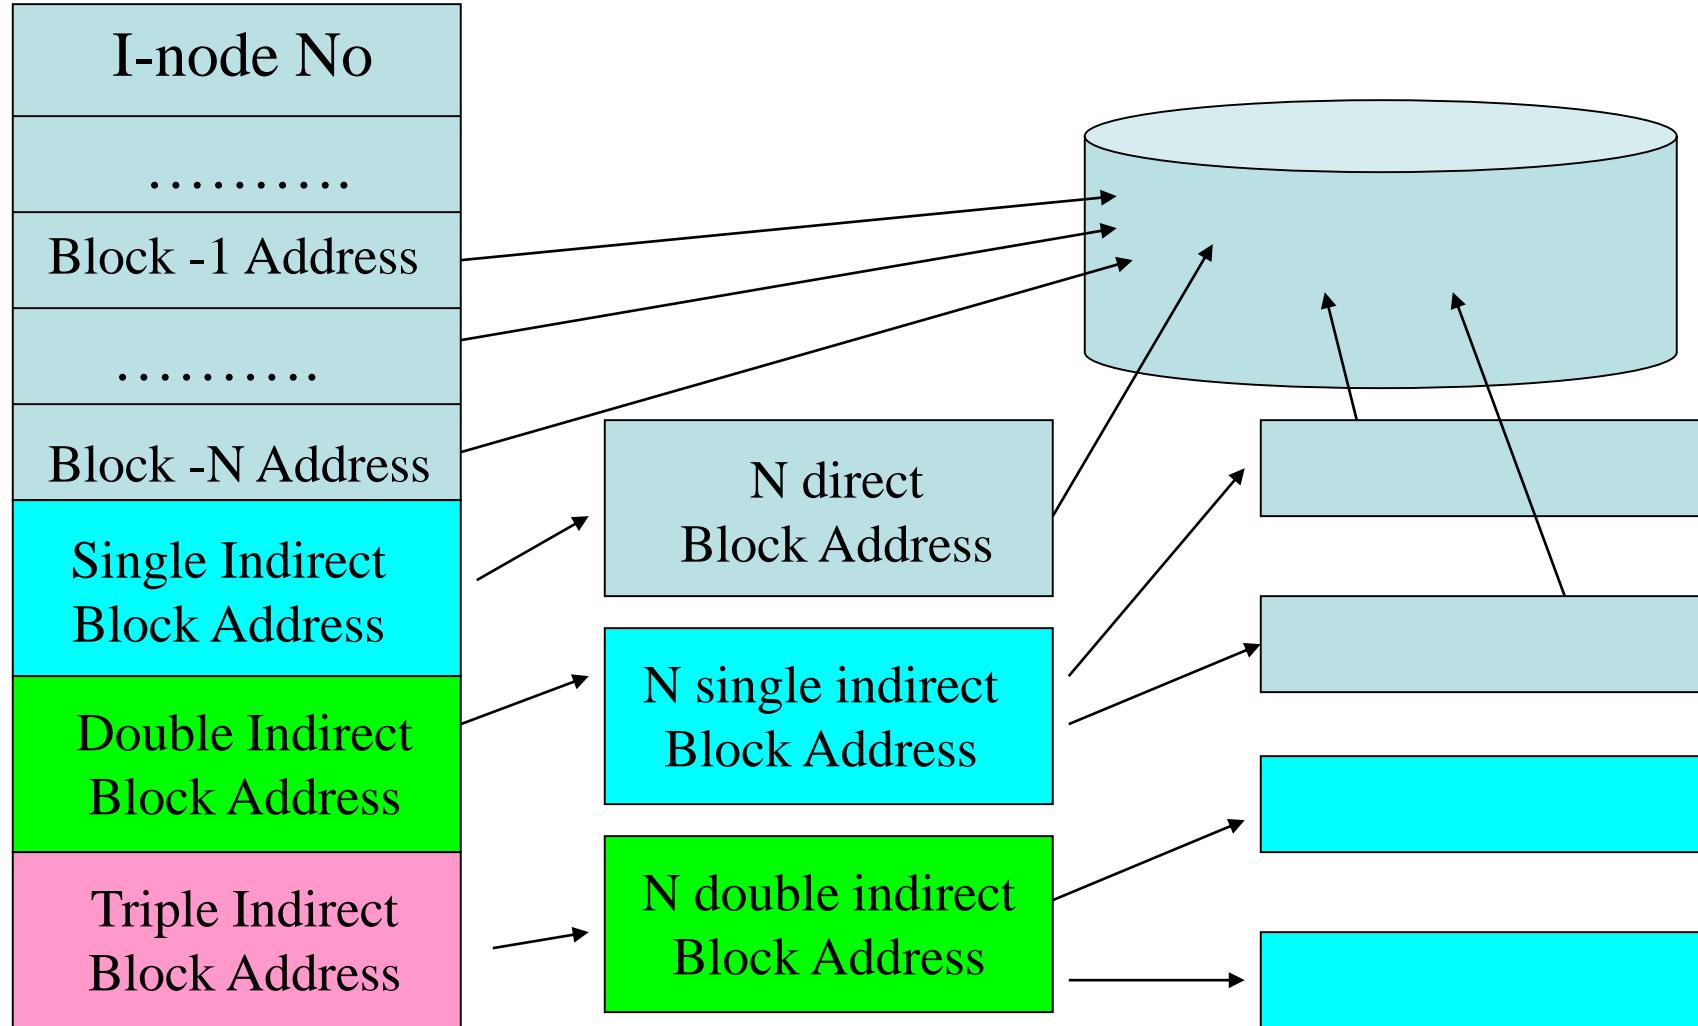
# File System Architecture



- Proc File System**
- 3 imp features
  - Process
  - System info
  - System Limitations



# ext File System



- Each I-node entry can track a very large file.
- Assuming a block size of 4KB and 12 direct data blocks, a single I-node entry can track file that has a maximum size of:
  - $12 * 4 \text{ KB} +$
  - $(1024 * 4 \text{ KB}) +$
  - $(1024 * 1024 * 4 \text{ KB}) +$
  - $(1024 * 1024 * 1024 * 4 \text{ KB})$
- File size can grow upto 16 terra bytes.

# XFS File System

When you create a file system, Linux creates a number of blocks on that device.



- Boot Block
  - Super-block
  - I-node table
  - Data Blocks
- Linux also creates an entry for the “/” (root) directory in the I-node table, and allocates data block to store the contents of the “/” directory.

• The super-block contains info. such as

- a bitmap of blocks on the device, each bit specifies whether a block is free or in use.
- the size of a data block
- the count of entries in the I-node table
- the date and time when the file system was last checked

# Process Management - Introduction

- Process is a program in execution.
- Processes carry out tasks in a system
- A process includes program counter (PC), CPU registers and process stacks, which contains temporary data.
- Linux is a multiprocessing system
- The Linux kernel is reentrant

- A process uses many resources like memory space, CPU, files, etc., during its lifetime.
- Kernel should keep track of the processes and the usage of system resources.
- Kernel should distributes resources among processes fairly.
- Most important resource is CPU. In a multiprocessing environment, to attain an ideal performance of a system, the CPU utilization should be maximum.

# Context Switch

- In order to run Unix, the computer hardware must provide two modes of execution
    - kernel mode
    - user mode
  - Some computers have more than two execution modes
    - eg: Intel processor. It has four modes of execution.
  - Each process has virtual address space , references to virtual memory are translated to physical memory locations using set of address translation maps.
- 
- Execution control is changing from one process to another.
  - When a current process either completes its execution or is waiting for a certain event to occur, the kernel saves the process context and removes the process from the running state.
  - Kernel loads next runnable process's registers with pointers for execution.
  - Kernel space: a fixed part of virtual address space of each process . It maps the kernel text and data structures.

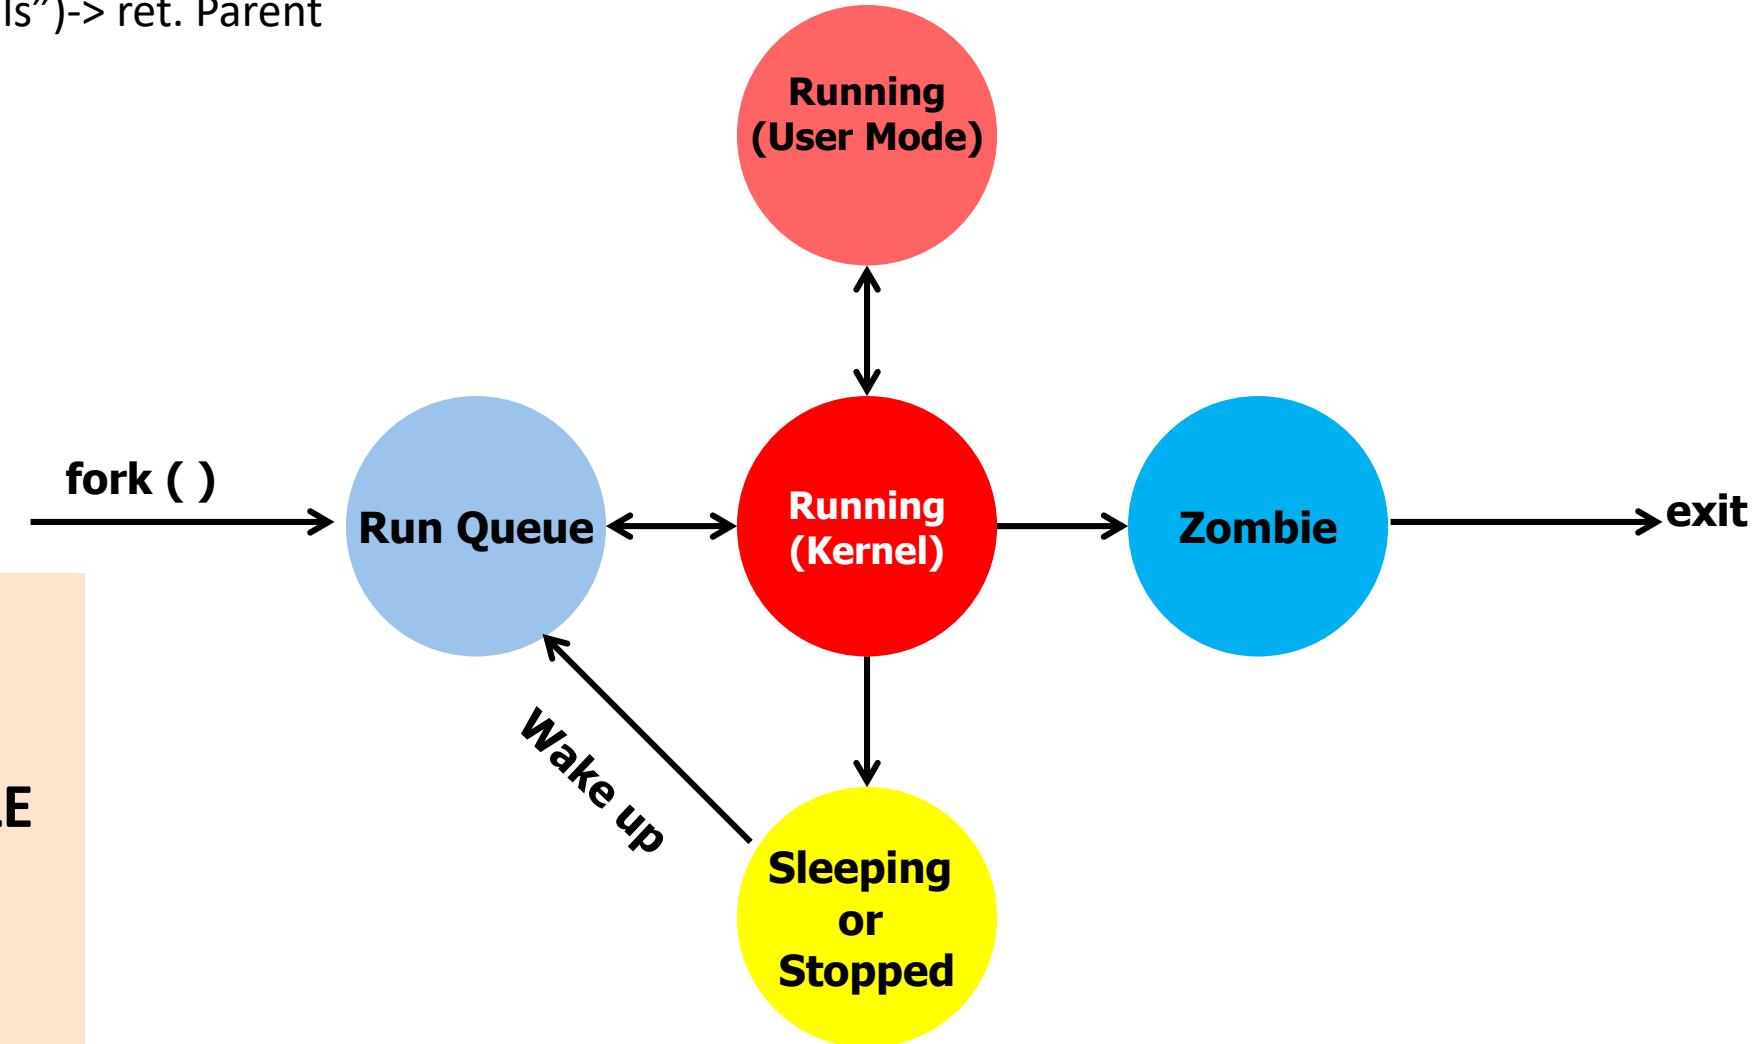
# Execution Context

- Kernel functions may execute either in process context or in system context
- User code runs in user mode and process context, and can access only the process space
- System calls and signals are handled in kernel mode but in process context, and may access process and system space
- Interrupts and system wide tasks are handled in kernel mode and system context, and must only access system space

- Every process is represented by a `task_struct` data structure.
- This structure is quite large and complex.
- Whenever a new process is created a new `task_struct` structure is created by the kernel and the complete process information is maintained by the structure.
- When a process is terminated, the corresponding structure is removed.
- Uses doubly linked list data structure.

# Process States

\$ls -> fork() -> child process-> exec("ls")-> ret. Parent



- **TASK\_RUNNING**
- **TASK\_INTERRUPTIBLE**
- **TASK\_UNINTERRUPTIBLE**
- **TASK\_STOPPED**
- **TASK\_ZOMBIE**

# Parent and Child

- **pid\_t fork (void); creates a new process**
- All statements after the `fork()` system call in a program are executed by two processes - the original process that used `fork()`, plus the new process that is created by `fork( )`

```
main ( ) {
    printf (" Hello fork %d\n, fork ( ) ");
}
```

- Hello fork: 0
- Hello fork: x (> 0);
- Hello fork: -1

```
if (!fork) {
    /* Child Code */
}
else {
    /* parent code */
    wait (0); /* or */
    waitpid(pid, ....);
}
```

Importance of “\n”.  
Process tree structure

# COW, Zombie and Orphan Process

- Instead of copying the address space of the parent, Linux uses the COW technique for economical use of the memory page.
- The parent space is not copied, it can be shared by both the parent and the child process but the memory pages are marked as write protected.
- if parent or child wants to modify the pages, then kernel copies the parent pages to the child process.
- Advantage: Kernel can defer or prevent copying of a parent process address space.

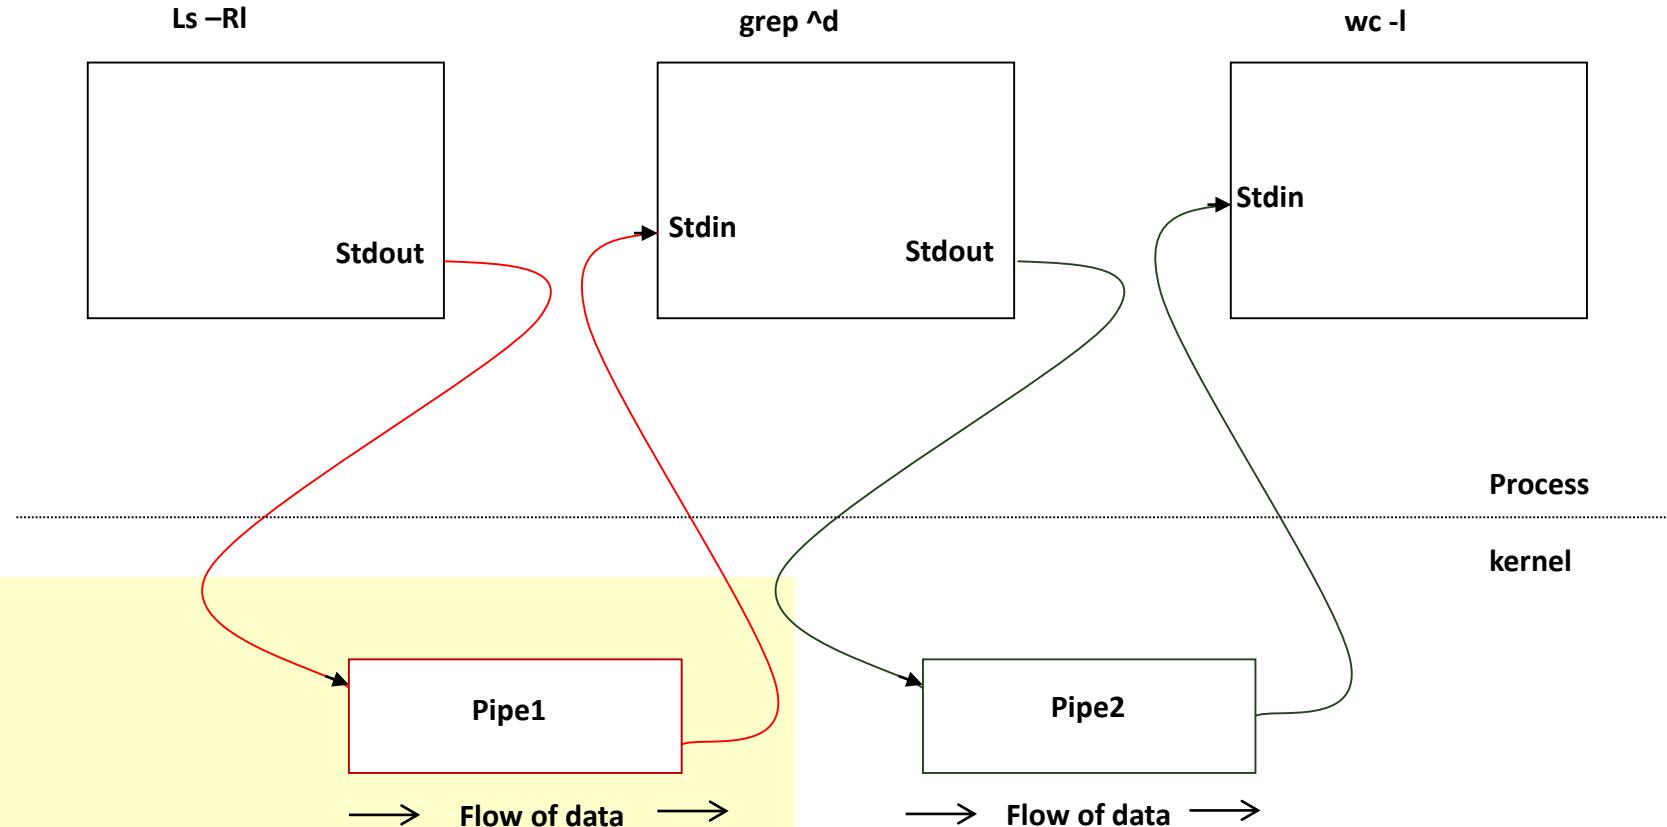
- When a child process exits, it has to give the exit status to the parent process.
- If the parent process is busy or suspended then the child process will not be able to terminate.
- Such state is called Zombie.
- if parent exits before child, the child will become an orphan process and the init process (grand parent) will take care of the child process.

# Inter Process Communications (IPC) Mechanisms

- In a multiprocessing environment, often many processes are in need to communicate with each other and share some of the resources.
- The shared resources must also be synchronized from the concurrent access by many processes.
- IPC mechanisms have many distinct purposes: for example
  - \* Data transfer
  - \* Event notification
  - \* Process control
  - \* Sharing data
  - \* Resource sharing

- Primitive
  - Unnamed pipe
  - Named pipe (FIFO)
- System V IPC
  - Message queues
  - Shared memory
  - Semaphores
- Socket Programming

# Execution of command: `$ ls -Rl | grep ^d | wc -l`



## Pipe advantages:

- Simplest form of IPC
- Persistence in process level
- Can be used in shell

## Disadvantages:

- Cannot be used to communicate between unrelated processes

# Execution of command: \$ ls -Rl | grep ^d | wc -l

```

int main()
{
    int fd1[2],fd2[2];

    pipe(fd1);
    pipe(fd2);

    if(!fork()) {
        dup2(fd1[1],1);
        close(fd1[0]); //closing read end of pipe
        close(fd2[0]);
        close(fd2[1]);
        execlp("ls","ls","-Rl",(char *)NULL);
    }
    else {
  
```

```

        if(!fork()) {
            dup2(fd1[0],0); //read end of first pipe
            dup2(fd2[1],1); //duplicated the write end of 2nd pipe
            close(fd1[1]); //closing the write end;
            close(fd2[0]); //closing read end of 2nd pipe
            execlp("grep","grep","^d",(char *)NULL);
        }
        else {
            dup2(fd2[0],0);
            close(fd2[1]);
            close(fd1[0]);
            close(fd1[1]);
            execlp("wc","wc","-l",(char *)NULL);
        }
    }
}
  
```

# Semaphore

- Synchronization Tool
- An Integer Number
- P ( ) And V ( ) Operators
- Avoid Busy Waiting
- Types of Semaphore

Used in :

- shared memory segment
- message queue
- file



# p and v operations

## Incrementing Operations:

```
int v (int i)
{
    i = i + 1; (unlock)
    return i;
}
```

- If a process wants to use the shared object, it will “lock” it by asking the semaphore to decrement the counter
- Depending upon the current value of the counter, the semaphore will either be able to carry out this operation, or will have to wait until the operation becomes possible
- The current value of counter is >0, the decrement operation will be possible. Otherwise, the process will have to wait

## Decrementing Operations:

```
int p (int i) {
    if (i > 0)
        then
            i--; (lock)
    else
        wait till i > 0;
    return i;
}
```

- System V semaphore provides a semaphore set - that can include a number of semaphores. It is up to user to decide the number of semaphores in the set
- Each semaphore in the set can be a binary or a counting semaphore. Each semaphore can be used to control access to one resource - by changing the value of semaphore count

# Semaphore Implementation

```

union semun {
    int val;      // value for SETVAL
    struct semid_ds *buf; // buffer for
    IPC_STAT, IPC_SET
    unsigned short int *array; // array for
    GETALL, SETALL
};

union semun arg;
semid = semget (key, 1, IPC_CREAT | 0644);
arg.val = 1;
/* 1 for binary else > 1 for Counting Semaphore */
semctl (semid, 0, SETVAL, arg);
  
```

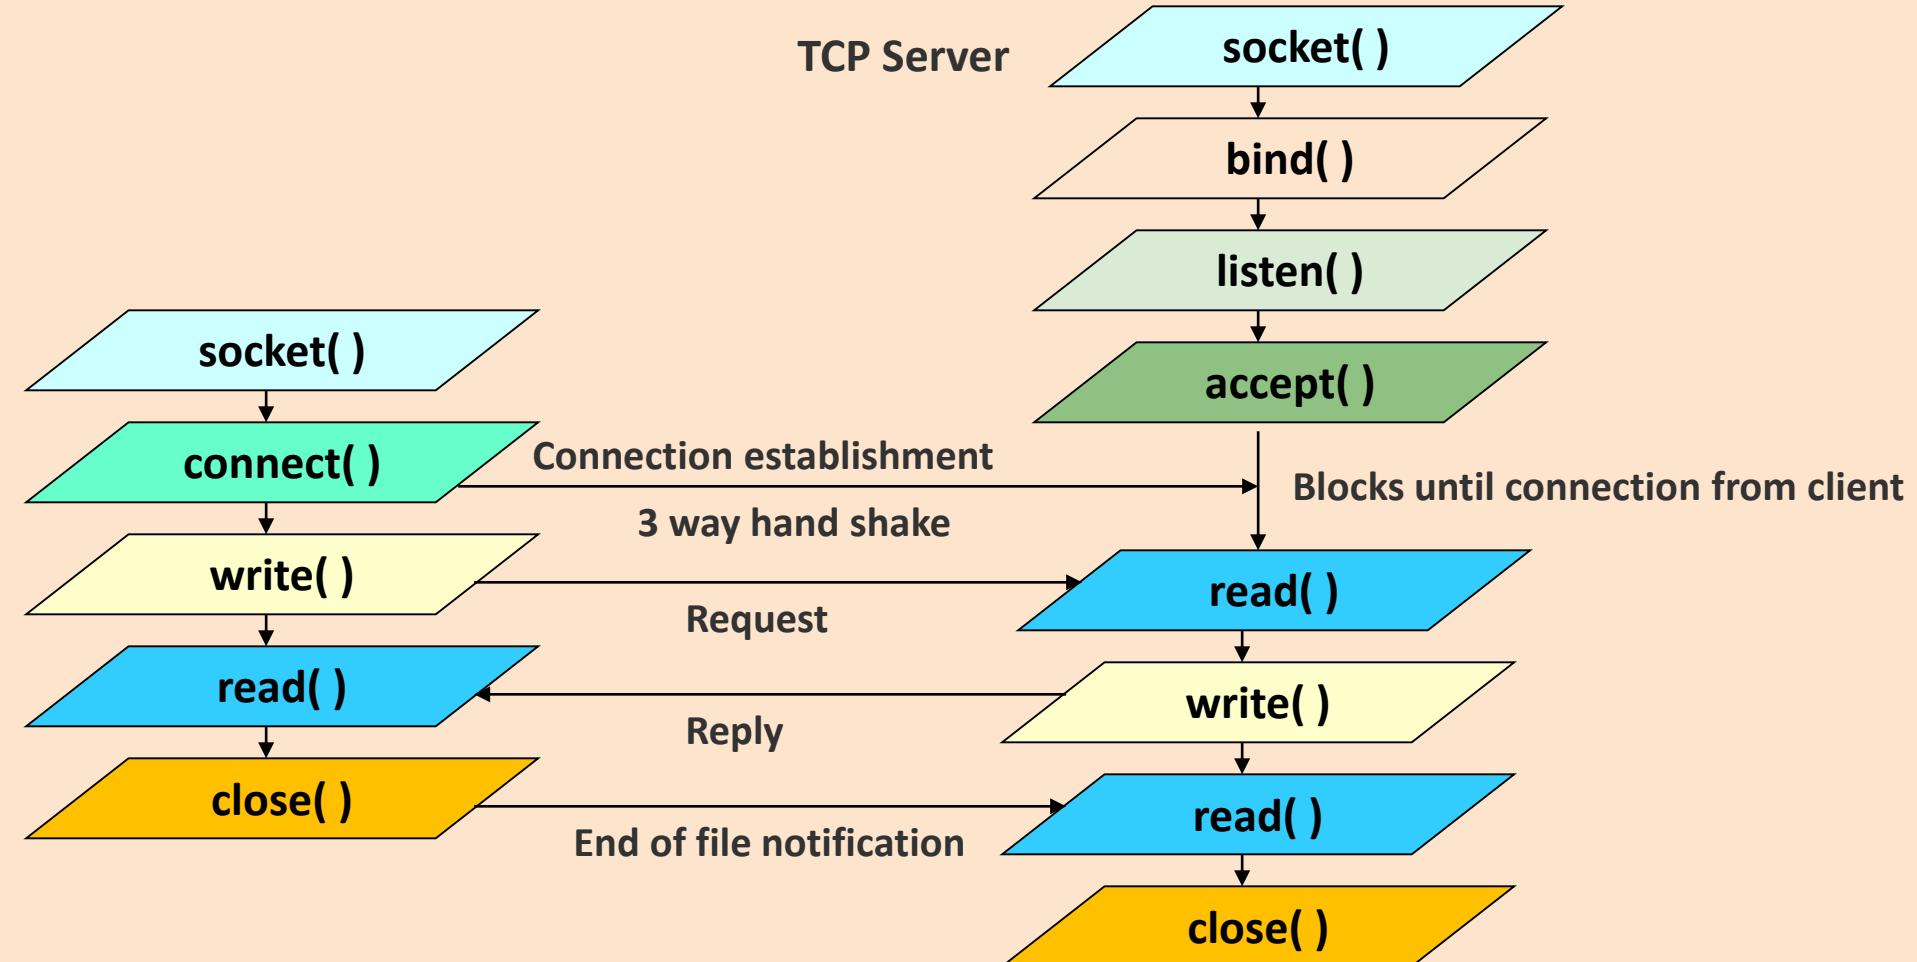
```

struct sembuf {
    short sem_num; /* semaphore number: 0 means first */
    short sem_op; /* semaphore operation: lock or unlock */
    short sem_flg; /* operation flags : 0, SEM_UNDO, IPC_NOWAIT */
};

struct sembuf buf = {0, -1, 0}; /* (-1 + previous value) */
semid = semget (key, 1, 0);

semop (semid, &buf, 1); /* locked */
-----Critical section-----
buf.sem_op = 1;
semop (semid, &buf, 1); /* unlocked */
  
```

# Socket Functions



# Alarm and Timers

- `unsigned int alarm (unsigned int seconds);`
- It is used to set an alarm for delivering **SIGALARM** signal.
- On success it returns zero.

- Three interval timers.
  - **ITIMER\_REAL**
    - This timer counts down in real (i.e., wall clock) time. At each expiration, a **SIGALRM** signal is generated.
  - **ITIMER\_VIRTUAL**
    - This timer counts down against the user-mode CPU time consumed by the process. (The measurement includes CPU time consumed by all threads in the process.) At each expiration, a **SIGVTALRM** signal is generated.
  - **ITIMER\_PROF**
    - This timer counts down against the total (i.e., both user and system) CPU time consumed by the process. At each expiration, a **SIGPROF** signal is generated.

# Time Stamp Counter

System can provide very high resolution time measurements through the time-stamp counter which counts the number of instructions since boot.

To measure Time Stamp Counter (TSC)

```
# include <sys/time.h>
unsigned long long rdtsc ( )
{
    unsigned long long dst;
    __asm__ __volatile__ ("rdtsc": "=A" (dst));
    return dst;
}
```

```
main ( )
{
    long long int start, end;
    start = rdtsc();

    /* Give your job; */

    end = rdtsc();

    printf (" Difference is : %llu\n", end - start);

}
/* This is the most accurate way of time measurement */
```

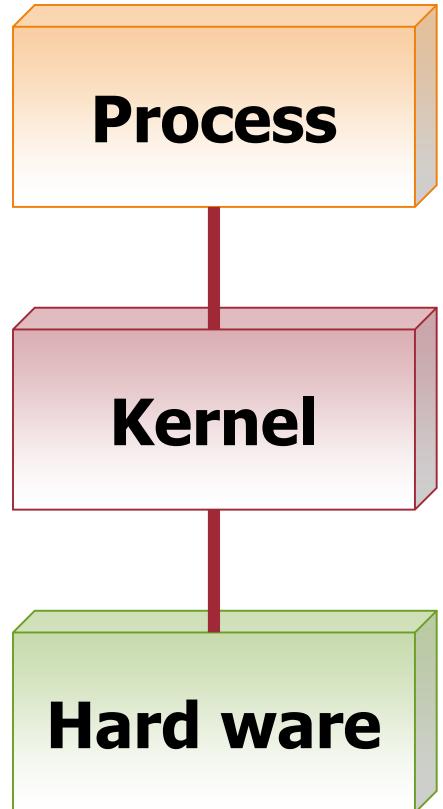
# Signals Introduction

- Signals are a fundamental method for inter process communication and are used in everything from network servers to media players.
- A signal is generated when
  - an event occurs (timer expires, alarm, etc.,)
  - a user quota exceeds (file size, no of processes etc.,)
  - an I/O device is ready
  - encountering an illegal instruction
  - a terminal interrupt like Ctrl-C or Ctrl-Z.
  - some other process send ( kill -9 pid)

- Each signal starts with macro SIGxxx.
- Each signal may also specifies with its integer number
- For help: \$ kill -l , \$ man 7 signal
- When a signal is sent to a process, kernel stops the execution and "forces" it to call the signal handler.
- When a process executes a signal handler, if some other signal arrives the new signal is blocked until the handler returns.

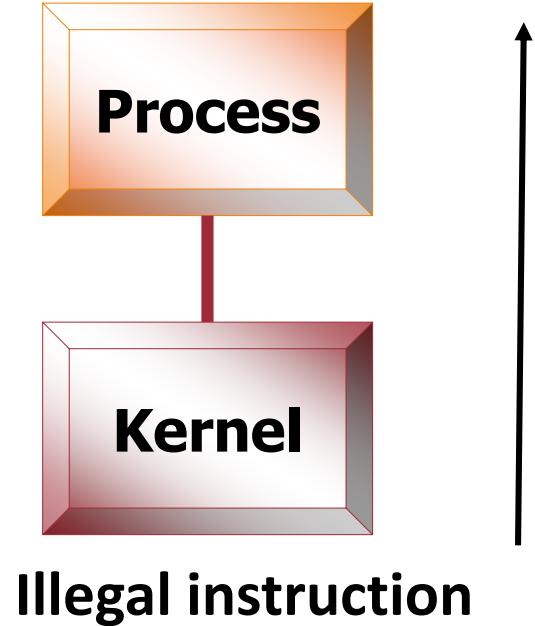
# Signal Vs Interrupt

## Interrupt



I/O operation (ex: mouse click)

## Signal



\$ kill -9 pid

# *signal* System Call

- How a process receives a signal, when it is
  - executing in user mode
  - executing in kernel mode
  - not running
  - in interruptible sleep state
  - in uninterruptible sleep state
- When a signal occurs, a process could
  - Catch the signal
  - Ignore the signal
  - Execute a default signal handler
- Two signals that cannot be caught or ignored
  - SIGSTOP
  - SIGKILL
- signal system call is used to catch, ignore or set the default action of a specified signal.
- `int signal (int signum, (void *) handler);`
- It takes two arguments: a signal number and a pointer to a user-defined signal handler.
- Two reserved predefined signal handlers are :
  - `SIG_IGN`
  - `SIG_DFL`